

Speech-To-Text and NLP: Comparing De Gasperi and Meloni's Rhetoric

MARCO BELLÒ, University of Padua, Italy

This manuscript compares Alcide De Gasperi and Giorgia Meloni rhetoric. After building a corpus of public speeches transcribed with Whisper from Giorgia Meloni's videos on YouTube, I compare them with a corpus of speeches from the Italian Republic's first Prime Minister Alcide De Gasperi. The results show that De Gasperi's speeches are, on average, 30% more diversified in their lexicon, and the corpora have a Jaccard similarity index of 32%. The highest-frequency words in common show transversal themes, such as freedom, peace and nation, while the highest-frequency words not in common reflect the PMs' different political and societal landscapes. I also performed a sentiment analysis on each speech using a BERT classifier. The results underscore the importance of domain-appropriate training data: the model was trained on product reviews, not on political rhetoric, which translated to an average confidence of roughly 40% per classification.

CCS Concepts: • **Computing methodologies** → **Natural language processing**; **Speech recognition**; • **Social and professional topics** → **Political speech**; • **Information systems** → **Sentiment analysis**.

Additional Key Words and Phrases: AI, NLP, Politics, Speeches

ACM Reference Format:

Marco Bellò. 2026. Speech-To-Text and NLP: Comparing De Gasperi and Meloni's Rhetoric. 1, 1 (May 2026), 10 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 INTRODUCTION

Politicians are some of the most influential people in human society: from the Greeks' demagogues to our Presidents and Prime Ministers, leaders have always tried to steer society. Shapiro and Page (1984) [13] showed how nearly all 20th century U.S. Presidents have tried to lead public opinion, with varying degrees of success. As societies became more democratic, so grew the leaders' responsibility to convince potential followers that their policies are beneficial. Thus, speeches play a vital role in the functioning of society [2]: leaders depend on the verbal power to persuade people [4].

Moreover, democratic stability depends on citizens on the losing side accepting election outcomes. Clayton et al. [6] evaluated the effects of exposure to multiple statements from (at the time) former president Donald Trump attacking the legitimacy of the 2020 US presidential election. They found no evidence indicating that elites can erode democratic norms easily or that the effects of norm violations are uniform across the entire population. However, elite rhetoric can shape normative beliefs in core democratic values such as confidence in elections and support for peaceful transfers of power. This was especially evident amongst people who approved of Trump's performance in office.

Having established the importance of speeches, we can differentiate between two types: *internal speeches*, aimed at other institutional bodies, such as parliamentary speeches, and *external speeches*, aimed at the population, such as public speeches, interviews and social network posts. This last category is especially interesting since it affords politicians a direct and high-volume communication channel to their potential followers: U.S. President Donald Trump tweeted

Author's address: Marco Bellò, marco.bello.vi@gmail.com, University of Padua, Padua, Italy.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

57,000 times between May 2009 and January 2021 alone [10], while Italy's Prime Minister Giorgia Meloni totalled more than eighty million social network interactions in 2025 [17].

On the other hand, internal speeches are especially easy to access, being already aggregated in public databases, yet their technical and codified nature makes their interpretation harder for both human and automated agents. Moreover, until 1985 Italian stenographic documentation did not report speeches word for word, but rather translated them into a "language of the parliament" [8, 12].

2 SOCIAL NETWORKS

Broadly speaking, the goal of this mini-project is to gather a corpus of politicians' speeches and perform linguistic analyses on it. The idea is to start from standard NLP tasks such as word frequency and text complexity, to then move to classification tasks such as sentiment analysis and hate speech detection.

This objective is why I ultimately decided to gather a corpus of *external speeches* rather than internal ones. Large language models are trained with tokenizers, and the resulting token distribution is highly imbalanced: Chung et al. [5] did a controlled study that scaled the vocabulary of the language model from 24K to 196K while holding data, computation, and optimization unchanged. They discovered that models are disproportionately optimized on the part of language that appears most often, meaning they are way better at understanding common words and patterns. This translates to the fact that internal speeches are less understandable for an AI agent.

As previously stated, external speeches include social network communications. While television is still the first source of information for the majority of the population, online platforms are steadily growing and are already the preferred media for young people [7, 11]. There is evidence of an increase in political participation due to social media usage, as well as risks for the functioning of democracy [1, 9], and there have been attempts to predict election results with social media data analyses. Rita et al. measure sentiment polarity on Twitter, and conclude that tweets' sentiment is not a reliable election results predictor. Additionally, results also show that it is impossible to state that social media impacts voting decisions [14].

Opposite results can be seen in Belcastro et al.'s [3] 2022 study: a real-time analysis was carried out during the 2020 US presidential election campaign, correctly identifying the leading candidate before Election Day in 10 out of 11 swing states.

Silva et al. [15] managed to match tweets against parliamentary speeches to measure politicians' sentiment on a specific topic. They find that parliament members who participate less in parliamentary debate tend to have larger differences with their party on Twitter, suggesting that a certain level of self-censoring is taking place in the parliamentary arena.

That being said, X is just one channel among many used to connect with the voters: official messages to the nation, interviews, and public speeches are still powerful tools, and are often made publicly available in video format by the politicians themselves on their Facebook page and YouTube channel.

3 PROJECT GOAL

Let's now define the goal of this project.

I want to compare the rhetoric between the first and the current¹ Prime Ministers of the Italian Republic, namely Alcide De Gasperi and Giorgia Meloni. The rhetoric will be compared using natural language processing techniques,

¹At the time of writing: May 2026.

such as measuring textual complexity and sentiment analysis. Only public speeches will be taken, thus leaving out all parliamentary speeches.

I also want to produce a proof-of-concept for a speech-to-text pipeline that, given a set of URLs, each of which is a YouTube's permalink, transcribes the video into text.

The first step will be gathering two corpora of public speeches, one for each PM. In the case of Giorgia Meloni, I want to compile the corpus by transcribing video interviews and public talks available on her official YouTube channel.

Then, after having gathered the corpora into suitable machine-readable data (such as csv), I will pre-process them with procedures such as tokenization, stemming and lemmatization. Lastly, I will compute word frequency and text complexity measures, as well as a sentiment analysis.

4 BUILDING THE CORPUS

4.1 Alcide De Gasperi's Corpus

I would like to thank Sara Tonelli for providing me with the De Gasperi's Corpus [16], a collection of Alcide De Gasperi's public documents with gold and silver annotation.

The corpus is a collection of 2,762 documents issued between 1901 and 1954, formatted into XML files which include metadata that covers not only the title, the date and the place of publication, but also key-concepts automatically extracted from each text (with the corresponding relevance score) and genre labels manually assigned by domain experts. Furthermore, the release includes silver annotation for lemma, part of speech, person names and place names with associated coordinates in a CoNLL-like format.

An example of the XML formatting can be seen at listing 1. Thanks to the `<genres>` tag, I was able to extract 474 public speeches, each of which with a location, a precise date and a list of keywords.

Listing 1. An example of how a single speech in De Gasperi's Corpus is formatted in XML

```

1 <?xml version="1.0" encoding="UTF-8" standalone="no"?>
2 <document id="I.doc_7">
3   <publication_date>1901-11-26</publication_date>
4   <publication_place>
5     <place_name>Trento</place_name>
6     <place_latitude>46.0664228</place_latitude>
7     <place_longitude>11.1257601</place_longitude>
8   </publication_place>
9   <keywords>
10    <keyword score="39.16">first keyword</keyword>
11    ...
12    <keyword score="5.59">last keyword</keyword>
13  </keywords>
14  <genres>
15    <genre>Speech Type</genre>
16  </genres>
17  <text>The speech itself if stored between these brackets</text>
18 </document>

```

4.2 Giorgia Meloni's Corpus

My objective was to build Meloni's corpus by transcribing videos of her public speeches, such as talks and interviews. Fortunately, there is an unofficial YouTube channel (that I reached from the Prime Minister's official website) which aggregates more than four thousand videos of her public appearances. The channel is called "Giorgia Meloni News"².

Given a YouTube URL, I can use Python's library *yt-dlp* to retrieve the video's metadata and download its audio content in mp3 format, as shown in listing 2. I had to manually pass the cookies taken from my web-browser, and I used some extra commands to prevent YouTube to block the requests due to suspicious activity.

Then, the mp3 file just retrieved gets injected into OpenAI's Whisper library, which uses an encoder-decoder Transformer to transcribe it into text, as shown in listing 3. The model I used is the "*medium*"³, which at 5 GB fits within the total 6 GB of VRAM available to my RTX 4050 GPU. The model was instructed to predict Italian.

The difference in precision between the models "tiny" (which requires less than 1 GB of VRAM) and "medium" is quite high, as can be seen in the following pieces of text that have been generated from the same audio file:

- **Model tiny:** "iamo con se è beruto fuori da questa rio neone sotto il profiro tecnico per quanto riguarda la franha e come il governo un tino era ad essere vicino alla popolazione di Nishenia."
- **Model medium:** "Diciamo cosa è venuto fuori da questa riunione sotto il profilo tecnico per quanto riguarda la frana e come il governo continuerà ad essere vicino alla popolazione diniscemeca."

The last step of the pipeline saves the transcription into a csv file, along with the extracted metadata. Each file has the following fields: *politician* ("meloni" in this case), *historical_date* (the upload date of the video), *location* and *tags* (extracted from the metadata if available, empty strings otherwise), *description* and *title* of the video, *url* which stores the permalink of the video itself, and lastly *text* which holds the whole transcription.

Each video is processed immediately after being retrieved, and until it has been saved in a csv file the program does not try to fetch the next one. This is done for two reasons: first, to prevent *yt-dlp* from sending requests too close to each other, which could result in YouTube denying them due to suspicious activity. Second, even if the pipeline halts for any reason, the videos processed up to that point will not be lost.

Listing 2. Downloading audio and metadata from a YouTube video

```
1 result = subprocess.run([
2     "yt-dlp",
3     "--print-json",
4     "-x",                                     # download audio only
5     "--audio-format", "mp3",
6     "--cookies", "yt_cookies.txt",          # manual cookies, extract with browser extension
7     "--sleep-requests", "2",                # Sleep 2s between requests
8     "--sleep-interval", "5",                # Sleep 5s between downloads
9     "--max-sleep-interval", "15",           # Randomize up to 15s
10    "--limit-rate", "5M",                    # Throttle to 5MB/s (mimics streaming)
11    "-o", audio_file.replace('.mp3', ''),
12    video_url
13 ], capture_output=True, text=True, check=True)
```

Listing 3. Transcribe the audio file into text using Whisper

```
1 subprocess.run([
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999
1000
```

²Giorgia Meloni News: <https://www.youtube.com/@GiorgiaMeloniTv>

³Whisper model sizes available are tiny, base, medium and large.

```

209 2  "whisper",
210 3  "--language",      lang_code,      # "Italian" chosen
211 4  "--word_timestamps", "True",
212 5  "--model",        model_name,    # model "medium" chosen
213 6  "--output_dir",    output_dir,
214 7  "--device",        "cuda",      # ensures the model gets loaded in GPU
215 8  audio_file
216 9 ], check=True)

```

5 TEXTUAL ANALYSIS

All the code used to compute the results shown in the next sections can be seen in the Jupyter notebook [mhetacc/cross-dem/notebooks/text_analysis.ipynb](https://github.com/mhetacc/cross-dem/notebooks/text_analysis.ipynb).

5.1 Token Count and Textual Complexity

The collected corpus contains 473 speeches from De Gasperi and 72 from Meloni, with a total token count of 609,142 for De Gasperi and 120,116 for Meloni. On average, Meloni's corpus has 1668 tokens per speech, while De Gasperi's has 1287 tokens per speech. Word metrics are shown in table 1. The text was tokenized with spaCy's Italian model *"it_core_news_sm"*.

I then measured the textual lexical diversity using *MTLD*, which I preferred over the more common *TTR* (Type-Token Ratio) since *MTLD* is agnostic to corpus size. I used the *LexicalRichness* library⁴ to compute the *MTLD* for each speech, and then both the average and the median *MTLD* for both PMs. Results can be seen in table 2.

To better visualize the difference in textual diversity, I plotted the *MTLD* values of 70 different speeches of Meloni and De Gasperi against each other. The plot can be seen in figure 1.

Table 1. Total tokens and average tokens per speech for De Gasperi and Meloni.

Politician	Total Tokens	Speech Count	Avg Tokens/Speech
De Gasperi	609,142	473	1,287.83
Meloni	120,116	72	1,668.28

Table 2. *MTLD* metrics for De Gasperi and Meloni.

Politician	Avg <i>MTLD</i>	Median <i>MTLD</i>
De Gasperi	141.830	137.642
Meloni	106.745	104.918

5.2 Vocabulary Comparison

I then wanted to see differences and similarities in the two Prime Ministers' vocabulary. I thus extracted from the corpus a filtered subset of tokens (example in listing 4), then I used the *Counter* method from the *collections* library to count the 1000 most commonly used words for both PMs.

⁴LexicalRichness is a small Python module to compute textual lexical richness (aka lexical diversity) measures: <https://pypi.org/project/lexicalrichness/>.




Fig. 1. Raw MTLD values for 70 speeches of Meloni and De Gasperi

Of the top-1000 words, 487 are in common, with a Jaccard similarity of 32.19%.

After extracting the top-1000 words by usage frequency for each PM, I cross-compared them. The results are shown in table 3. The first column shows words that are in common, whose compound frequency (*meloni_freq[word] + degasperi_freq[word]*) is highest. The second column shows words that are only found in De Gasperi's corpus (Meloni's frequency for those words is zero). The third column shows words that are only found in Meloni's corpus (De Gasperi's frequency for those words is zero).

It is interesting to see that words in common are about transversal themes, such as *italia*, *popolo*, *libertà*, and *democrazia*. Interestingly, *europa* is a top-20 shared word, showing us that the theme was and still is relevant even though the two PMs were born almost 100 years apart from each other.

On the other hand, the PM-specific words show us how the public discourse changes in the different periods: there are some themes that seem to have disappeared; for example, in the De Gasperi-only top 20 we see words such as *comunismo*, *socialisti*, and *marshall*, all of which were extremely relevant 70 years ago. In the Meloni-only top 20 we find words that simply could not have mattered back then, even though today they are absolutely central, such as *trump*, *ucraina*, *putin*, and *hamas*.

We can also see some examples of how the language itself changes over time: De Gasperi used quite often the word *dinanzi*, while Meloni often uses *dopodiché*. Both of these words are never used by the other Prime Minister.

To try to visualize how close (or far apart) the two PM's vocabularies are, I visualized their top-1000 words in a t-SNE scatter plot using the *scikit-learn* library. The results (in figure 2) show us that, with the exception of the words in common (in purple), the two vocabularies seem to be quite far apart, clustering in different regions.

Listing 4. Extract and filter words with tokenize function.

```
def tokenize(text):
    return [token.text.lower() for token in nlp(text)
            if token.is_alpha      # only alphanumeric
            and not token.is_stop  # no stop-words
            and len(token)>3       # removes articles and some filler words
            ]
```

Table 3. Comparison of most frequent words common to both PM. Last two columns show words specific to only one PM.

Rank	Common Words (Sum)		De Gasperi Only (Freq)		Meloni Only (Freq)	
1	italia	(1967)	gasperi	(458)	euro	(88)
2	popolo	(1201)	comunisti	(412)	ucraina	(85)
3	libertà	(1071)	trieste	(309)	sottotitoli	(65)
4	guerra	(927)	comunista	(241)	dopodiché	(48)
5	politica	(912)	avvenire	(237)	trump	(46)
6	pace	(910)	togliatti	(207)	infrastrutture	(37)
7	italiano	(828)	trento	(173)	approccio	(32)
8	presidente	(733)	comunismo	(171)	competitività	(30)
9	democrazia	(708)	socialisti	(170)	gaza	(28)
10	partito	(682)	oratore	(165)	migrazione	(28)
11	bisogna	(661)	deputati	(158)	scenario	(26)
12	italiani	(605)	trentino	(149)	centrodestra	(26)
13	soprattutto	(567)	dittatura	(145)	record	(23)
14	europa	(563)	democratici	(143)	hamas	(23)
15	sociale	(548)	nenni	(142)	palestina	(20)
16	cioè	(535)	voti	(137)	meloni	(19)
17	altra	(530)	degasperi	(127)	food	(19)
18	amici	(522)	onorevole	(126)	mattoni	(19)
19	nazione	(508)	dinanzi	(125)	buongiorno	(19)
20	senso	(502)	marshall	(122)	putin	(19)




Fig. 2. t-SNE for top-1000 words for both PMs. In red Meloni-only words, in blue De Gasperi-only words, and in purple shared words.

5.3 Critiques

A better filtering of the extracted words would be needed: for example, "gasperi" and "degasperi" should ideally be merged into the bigram "de gasperi", if not removed entirely, since he did not refer to himself in the third person.

This could be achieved by applying named entity recognition (NER) to identify and merge multi-word names, or by post-processing the token list to combine frequent bigrams corresponding to named entities.

Additionally, a custom stop-word list could be compiled: for example, in Meloni's corpus the word "sottotitoli" appears often, but was likely added in post-production by the video uploader.

6 SENTIMENT ANALYSIS

The last analysis I performed on the dataset was a sentiment analysis using the *transformers* library from Hugging Face.

Originally I wanted to use *MilaNLProc/feel-it-italian-emotion*⁵, a model trained on the FEEL-IT corpus, made of Italian Twitter posts annotated with four basic emotions: anger, fear, joy, sadness. Unfortunately, at the time of writing the model's tokenizer is not compatible with any Python version I tried.

I opted for one of the most popular free models, *nlptown/bert-base-multilingual-uncased-sentiment*⁶, a bert-base-multilingual-uncased model fine-tuned for sentiment analysis on product reviews in six languages: English, Dutch, German, French, Spanish, and Italian. It predicts the sentiment as a number of stars (between 1 and 5).

I used a *pipeline* from the *transformers* library, and I ran the classifier on 512 tokens at a time, which is BERT's maximum input length.

The classifier returns two results: a label (between 1 and 5, as in "4 stars"), and a score, which represents how confident the model is about the label itself. For example, if a piece of text is labelled as four stars with a score of 0.5, the model is only 50% sure that that text should be labelled as four stars.

To make comparison easier, I aggregated labels and scores with the following formula:

$$\frac{\text{amount of stars} \times \text{score}}{5}$$

Lastly, I appended the lists with all the scores to the original datasets.

A code example can be seen in listing 5. De Gasperi's median aggregated sentiment is 0.265, while Meloni's median aggregated sentiment is 0.379. A visualization of the aggregated sentiment can be seen in figure 3. The samples are picked at random.

Listing 5. Example code for a function that calculates the sentiment score for each piece of text.

```

1 sentiment_classifier = pipeline(
2     "text-classification",
3     model="nlptown/bert-base-multilingual-uncased-sentiment",
4     device=0 # GPU
5 )
6
7 sentiment_results = sentiment_classifier(
8     texts, # list with all speeches
9     batch_size=batch_size,
10    truncation=True,
11    max_length=512 # BERT max input length
12 )
13

```

⁵<https://huggingface.co/MilaNLProc/feel-it-italian-emotion>

⁶<https://huggingface.co/nlptown/bert-base-multilingual-uncased-sentiment>.


```

417 14 df['sentiment_aggregated'] = [
418 15     (int(res['label'].split()[0]) * res['score']) / 5
419 16     for res in sentiment_results
420 17 ]
421 18 return df

```




Fig. 3. Aggregated sentiment for 70 randomly sampled speeches. Meloni's speeches in red, De Gasperi's speeches in blue.

6.1 Critiques and Future Work

The model used to perform the sentiment analysis was designed to classify product reviews, and it struggled with long political speeches. This is evident from the sentiment scores (shown in table 4), where the model's confidence in its classifications is around 40%.

It would be interesting to perform a sentiment analysis on the same corpus with a model trained on, for example, newspaper articles.

In future work, fine-tuning a model on labelled political speeches should produce better classification output.

Table 4. Mean and median sentiment scores of the sentiment analysis for Meloni and De Gasperi.

Metric	Meloni	De Gasperi
Mean Sentiment Score	0.4636	0.4068
Median Sentiment Score	0.4692	0.3627

7 CONCLUSION

Overall, the results are satisfactory. I managed to prototype a pipeline that retrieves YouTube videos, transcribes them in text with an encoder-decoder transformer, and saves the result in a csv file. I retrieved 72 speeches of Giorgia Meloni, for a total of 120, 116 tokens.

I compared them with the De Gasperi corpus, after extracting 473 public speeches for a total of 609, 142 tokens. I calculated the MTLT for both corpora, showing that De Gasperi's speeches are, on average, 30% more diversified in their lexicon.

Then, I compared the two Prime Ministers' 1000 most-used words: they have 487 words in common, with a Jaccard similarity index of 32%. The highest-frequency words in common show transversal themes, such as freedom, peace and

nation, while the highest-frequency words not in common reflect the PM's different political and societal landscapes. De Gasperi often talked about communism and socialism, while Meloni focuses on Ukraine, Trump and the Euro.

Lastly, I performed a sentiment analysis on each speech using a BERT classifier. The results underscore the importance of domain-appropriate training data: the model was trained on product reviews, not on political rhetoric, which translated to an average confidence of roughly 40% per classification.

REFERENCES

- [1] Eran Amsalem and Alon Zoizner. 2023. Do people learn about politics on social media? A meta-analysis of 76 studies. *Journal of Communication* 73, 1 (Feb. 2023), 3–13. <https://doi.org/10.1093/joc/jqac034>
- [2] Adrian Beard. 2000. *The Language of Politics*. Routledge. <https://newuniversityinexileconsortium.org/wp-content/uploads/2022/02/the-language-of-politics.pdf> Google-Books-ID: XpqSmh1pjsEC.
- [3] Loris Belcastro, Francesco Branda, Riccardo Cantini, Fabrizio Marozzo, Domenico Talia, and Paolo Trunfio. 2022. Analyzing voter behavior on social media during the 2020 US presidential election campaign. *Social Network Analysis and Mining* 12, 1 (July 2022), 83. <https://doi.org/10.1007/s13278-022-00913-9>
- [4] Jonathan Charteris-Black. 2011. Persuasion, Speech Making and Rhetoric. In *Politicians and Rhetoric: The Persuasive Power of Metaphor*, Jonathan Charteris-Black (Ed.). Palgrave Macmillan UK, London, 1–27. https://doi.org/10.1057/9780230319899_1
- [5] Woojin Chung and Jeonghoon Kim. 2025. Exploiting Vocabulary Frequency Imbalance in Language Model Pre-training. <https://doi.org/10.48550/arXiv.2508.15390> arXiv:2508.15390 [cs].
- [6] Katherine Clayton, Nicholas T. Davis, Brendan Nyhan, Ethan Porter, Timothy J. Ryan, and Thomas J. Wood. 2021. Elite rhetoric can undermine democratic norms. *Proceedings of the National Academy of Sciences* 118, 23 (June 2021), e2024125118. <https://doi.org/10.1073/pnas.2024125118>
- [7] Eurobarometer. 2023. *Media & News Survey 2023 - November 2023 - - Eurobarometer survey*. Technical Report. European Union. <https://doi.org/10.2861/11595>
- [8] Günter Holtus and Michele Cortelazzo. 1985. Dal parlato al (tra)scritto: i resoconti stenografici dei discorsi parlamentari. In *Gesprochenes Italienisch in Geschichte und Gegenwart*. G. Narr, 86–118. https://www.uni-saarland.de/fileadmin/upload/fachrichtung/romanistik/W._Schweickard/Publikationen/Gesprochenes_Italienisch_in_Geschichte_und_Gegenwart.pdf Google-Books-ID: NfwlAAAAMAAJ.
- [9] Philipp Lorenz-Spreen, Lisa Oswald, Stephan Lewandowsky, and Ralph Hertwig. 2023. A systematic review of worldwide causal and correlational evidence on digital media and democracy. *Nature Human Behaviour* 7, 1 (Jan. 2023), 74–101. <https://doi.org/10.1038/s41562-022-01460-1>
- [10] Aamer Madhani and Jill Colvin. 2021. A farewell to @realDonaldTrump, gone after 57,000 tweets. <https://apnews.com/article/twitter-donald-trump-ban-cea450b1f12f4ceb8984972a120018d5>
- [11] Bron Maher. 2024. Twice as many Brits got 2024 election news from TV as from social media, Ofcom finds. <https://pressgazette.co.uk/news/ofcom-2024-election-news-sources-television-social-media-working-class/>
- [12] Aurelia Mohrhoff. 1987. Dalla lingua del Parlamento alla lingua del parlamentare. In *Aurelia Mohrhoff, Dalla lingua del Parlamento alla lingua del parlamentare*. Roma: Camera dei Deputati, 145–156. https://bpr.camera.it/bpr/allegati/show/3815_828_t
- [13] Benjamin I. Page and Robert Y. Shapiro. 1984. Presidents as Opinion Leaders: Some New Evidence. *Policy Studies Journal* 12, 4 (June 1984), 649–662. <https://doi.org/10.1111/j.1541-0072.1984.tb00480.x>
- [14] Paulo Rita, Nuno António, and Ana Patrícia Afonso. 2023. Social media discourse and voting decisions influence: sentiment analysis in tweets during an electoral period. *Social Network Analysis and Mining* 13, 1 (March 2023), 46. <https://doi.org/10.1007/s13278-023-01048-1>
- [15] Bruno Castanho Silva and Sven-Oliver Proksch. 2022. Politicians unleashed? Political communication on Twitter and in parliament in Western Europe. *Political Science Research and Methods* 10, 4 (Oct. 2022), 776–792. <https://doi.org/10.1017/psrm.2021.36>
- [16] Sara Tonelli, Rachele Sprugnoli, and Giovanni Moretti. 2019. Prendo la Parola in Questo Consesso Mondiale:. (2019). <https://ceur-ws.org/Vol-2481/paper71.pdf>
- [17] Alessandro Vinci. 2025. Meloni leader anche sui social, 3,3 milioni di nuovi follower nel 2025: staccato Salvini. Schlein prima per engagement. https://www.corriere.it/tecnologia/25_dicembre_17/meloni-leader-anche-sui-social-3-3-milioni-di-nuovi-follower-nel-2025-staccato-salvini-schlein-prima-per-engagement-431d8cda-a6e3-4d04-880b-922639075xlk.shtml Section: Meloni leader anche sui social, 3,3 milioni di nuovi follower nel 2025: staccato Salvini. Schlein prima per engagement.

Received 26/05/2026